

# CS5222 Project 2

## Custom Acceleration with FPGAs

Shen Jiamin  
A0209166A  
shen\_jiamin@u.nus.edu

March 5, 2022

### Abstract

In this project, I'm going to port the lab to **PYNQ 2.7** and **Vivado/Vitis 2020.2**. The experiment is done on ASUS RS500-E8-PS4 V2, with operating system Ubuntu 20.04.4 LTS (GNU/Linux 5.4.0-100-generic x86\_64).

## 1 Matrix Multiplication Pipeline Optimization in HLS

### 1.1 Understanding the baseline matrix multiply (background)

For Vitis 2020.2, the command used should be

```
$ vitis_hls -f hls.tcl
```

The report generated by HLS (as in [Table 1](#)) shows that some pipelining has already been done automatically by Vitis HLS. In order to prepare baseline for the next part, I disabled the pipelining.

```
--- hls.tcl          2022-03-03 21:17:24.651417872 +0800
+++ hls_nopipe.tcl   2022-03-03 21:33:53.435003340 +0800
@@ -7,6 +7,7 @@
open_solution "solution0" -flow_target vivado
set_part {xc7z020clg484-1}
create_clock -period 10 -name default
+config_compile -pipeline_loops 0
csim_design -clean
csynth_design
close_project
```

The new report is as [Table 2](#). It turns out that the overall performance is a little bit worse than documented. This is because every iteration in L3 loop takes 11 cycles and thus 2816 cycles in total to perform a single inner product.

Table 1: Loop details for baseline with automatic pipelining

| Loop Name             | Latency (cycles) |       | Iteration<br>Latency | Initiation Interval |        | Trip<br>Count | Pipelined |
|-----------------------|------------------|-------|----------------------|---------------------|--------|---------------|-----------|
|                       | min              | max   |                      | achieved            | target |               |           |
| - LOAD_OFF_1          | 5                | 5     | 1                    | 1                   | 1      | 5             | yes       |
| - LOAD_W_1_LOAD_W_2   | 1280             | 1280  | 2                    | 1                   | 1      | 1280          | yes       |
| - LOAD_I_1_LOAD_I_2   | 1024             | 1024  | 2                    | 1                   | 1      | 1024          | yes       |
| - L1_L2               | 82800            | 82800 | 1035                 | -                   | -      | 80            | no        |
| + L3                  | 1031             | 1031  | 12                   | 4                   | 1      | 256           | yes       |
| - STORE_O_1_STORE_O_2 | 42               | 42    | 4                    | 1                   | 1      | 40            | yes       |

Table 2: Loop details for baseline without automatic pipelining

| Loop Name    | Latency (cycles) |        | Iteration Latency | Initiation Interval |        | Trip Count | Pipelined |
|--------------|------------------|--------|-------------------|---------------------|--------|------------|-----------|
|              | min              | max    |                   | achieved            | target |            |           |
| - LOAD_OFF_1 | 5                | 5      | 1                 | -                   | -      | 5          | no        |
| - LOAD_W_1   | 1300             | 1300   | 130               | -                   | -      | 10         | no        |
| + LOAD_W_2   | 128              | 128    | 1                 | -                   | -      | 128        | no        |
| - LOAD_I_1   | 1040             | 1040   | 130               | -                   | -      | 8          | no        |
| + LOAD_I_2   | 128              | 128    | 1                 | -                   | -      | 128        | no        |
| - L1         | 225536           | 225536 | 28192             | -                   | -      | 8          | no        |
| + L2         | 28190            | 28190  | 2819              | -                   | -      | 10         | no        |
| ++ L3        | 2816             | 2816   | 11                | -                   | -      | 256        | no        |
| - STORE_O_1  | 136              | 136    | 17                | -                   | -      | 8          | no        |
| + STORE_O_2  | 15               | 15     | 3                 | -                   | -      | 5          | no        |

Table 3: Performance and utilization estimates for mmult\_float

| Profile |                     | Latency (cycles) |        | Latency (ms) |       | Interval (cycles) |        | Pipeline Type |
|---------|---------------------|------------------|--------|--------------|-------|-------------------|--------|---------------|
|         |                     | min              | max    | min          | max   | min               | max    |               |
| 1.1     | Baseline (AutoPipe) | 85160            | 85160  | 1.236        | 1.236 | 85161             | 85161  | none          |
| 1.1     | Baseline (NoPipe)   | 228022           | 228022 | 2.280        | 2.280 | 228023            | 228023 | none          |
| 1.2.1   | L3 Pipelining       | 85286            | 85286  | 1.238        | 1.238 | 85287             | 85287  | none          |
| 1.2.3   | L1 Pipelining       | 6193             | 6193   | 0.062        | 0.062 | 6194              | 6194   | none          |
| 1.2.2   | L2 Pipelining       | 7341             | 7341   | 0.073        | 0.073 | 7342              | 7342   | none          |

  

| Profile |                     | Utilization Summary |     |        |        |      | Instance |      |
|---------|---------------------|---------------------|-----|--------|--------|------|----------|------|
|         |                     | BRAM_18K            | DSP | FF     | LUT    | URAM | fadd     | fmul |
|         | Available           | 280                 | 220 | 106400 | 53200  | 0    |          |      |
| 1.1     | Baseline (AutoPipe) | 13                  | 5   | 1151   | 2058   | 0    | 1        | 1    |
| 1.1     | Baseline (NoPipe)   | 14                  | 5   | 817    | 1635   | 0    | 1        | 1    |
| 1.2.1   | L3 Pipelining       | 14                  | 5   | 921    | 1713   | 0    | 1        | 1    |
| 1.2.3   | L1 Pipelining       | 70                  | 800 | 415044 | 243128 | 0    | 160      | 160  |
| 1.2.2   | L2 Pipelining       | 182                 | 80  | 38357  | 34359  | 0    | 16       | 16   |

## 1.2 Pipelining in HLS (8 marks)

The work is done with auto pipelining disabled. The performance and utilization estimates are reported in Table 3.

### 1.2.1 Pipelining the L3 (innermost) loop

The code is modified as Figure 1. As reported in Table 4, pipelining L3 reduces its latency from 2816 cycles to 1031 cycles. The L1 is flattened, but L2 cannot be flattened because L2 is not a perfect loop. Thus the result has 2-layer loops where the outer layer has 80 iterations, and the inner layer has 256 iterations. There is no parallelism in this condition. This design utilizes slightly more resources but no more floating point adders or multipliers. The overall latency is 85285 cycles, which is about 2.67x speedup. Other statistics in detail can be found in Table 3.

### 1.2.2 Pipelining the L2 loop

The code is modified as Figure 2. As reported in Table 5, pipelining the loop body of L2 unrolls L3 and introduces both pipelining and some parallelism, which reduces the iteration latency from 2816 cycles to 1287 cycles and the total latency for the matrix multiplication from 225536 to 2550. The design used 16 adders and 16 multipliers. The overall latency is 7341 cycles, about 31.06x speedup relative to baseline. Other statistics in detail can be found in Table 3.

```

--- mmult_float.cpp.orig      2022-03-04 16:59:04.734375380 +0800
+++ mmult_float.cpp.L3       2022-03-04 16:58:59.566207094 +0800
@@ -78,6 +78,7 @@
     T tmp = offset_buf[j];
L3:
    for (int k = 0; k < FEAT; k++) {
+#pragma HLS PIPELINE II = 1
        tmp += in_buf[i][k] * weight_buf[j][k];
    }
    out_buf[i][j] = tmp;

```

Figure 1: Inserting HLS directive for L3.

Table 4: Loop details for L3 pipelining

| Loop Name    | Latency (cycles) |       | Iteration<br>Latency | Initiation Interval |        | Trip<br>Count | Pipelined |
|--------------|------------------|-------|----------------------|---------------------|--------|---------------|-----------|
|              | min              | max   |                      | achieved            | target |               |           |
| - LOAD_OFF_1 | 5                | 5     | 1                    | -                   | -      | 5             | no        |
| - LOAD_W_1   | 1300             | 1300  | 130                  | -                   | -      | 10            | no        |
| + LOAD_W_2   | 128              | 128   | 1                    | -                   | -      | 128           | no        |
| - LOAD_I_1   | 1040             | 1040  | 130                  | -                   | -      | 8             | no        |
| + LOAD_I_2   | 128              | 128   | 1                    | -                   | -      | 128           | no        |
| - L1_L2      | 82800            | 82800 | 1035                 | -                   | -      | 80            | no        |
| + L3         | 1031             | 1031  | 12                   | 4                   | 1      | 256           | yes       |
| - STORE_O_1  | 136              | 136   | 17                   | -                   | -      | 8             | no        |
| + STORE_O_2  | 15               | 15    | 3                    | -                   | -      | 5             | no        |

```

--- mmult_float.cpp.orig      2022-03-04 16:59:04.734375380 +0800
+++ mmult_float.cpp.L2       2022-03-04 17:00:26.201027719 +0800
@@ -74,6 +74,7 @@
// Iterate over output classes
L2:
    for (int j = 0; j < CLASSES; j++) {
+#pragma HLS PIPELINE II = 1
        // Perform the dot product
        T tmp = offset_buf[j];
L3:

```

Figure 2: Inserting HLS directive for L2.

Table 5: Loop details for L2 pipelining

| Loop Name    | Latency (cycles) |      | Iteration<br>Latency | Initiation Interval |        | Trip<br>Count | Pipelined |
|--------------|------------------|------|----------------------|---------------------|--------|---------------|-----------|
|              | min              | max  |                      | achieved            | target |               |           |
| - LOAD_OFF_1 | 5                | 5    | 1                    | -                   | -      | 5             | no        |
| - LOAD_W_1   | 2580             | 2580 | 258                  | -                   | -      | 10            | no        |
| + LOAD_W_2   | 256              | 256  | 2                    | -                   | -      | 128           | no        |
| - LOAD_I_1   | 2064             | 2064 | 258                  | -                   | -      | 8             | no        |
| + LOAD_I_2   | 256              | 256  | 2                    | -                   | -      | 128           | no        |
| - L1_L2      | 2550             | 2550 | 1287                 | 16                  | 1      | 80            | yes       |
| - STORE_O_1  | 136              | 136  | 17                   | -                   | -      | 8             | no        |
| + STORE_O_2  | 15               | 15   | 3                    | -                   | -      | 5             | no        |

### 1.2.3 Pipelining the L1 (outermost) loop

The code is modified as Figure 3. The loop details are in Table 6. Pipelining L1 makes both L2 and L3 completely unrolled, which makes there only one loop with 8 iterations. The unrolled loop body is heavily paralleled, which make use of 160 floating point adders and 160 floating point multipliers. The parallelism reduces the latency for one iteration from 28190 cycles to 1291 cycles. The pipelining further reduce the latency for L1 to 1402 cycles, although there are eight iterations. The overall latency is 6193 cycles, about 36.82x speedup relative to L3 pipelining and 1.19x speedup compared to L2 pipelining. Other statistics in detail can be found in Table 3.

Although L1 pipelining achieves some speedup compared to L2 pipelining, it takes 299.1 seconds to complete the whole building process, while the time for L2 pipelining is only 62.5 seconds. At the same time, the hardware resource usage has exceeded those available on board.<sup>1</sup> Thus, L2 pipelining will be a better result after the tradeoff between performance and resource usage.

```
--- mmult_float.cpp.orig      2022-03-04 16:59:04.734375380 +0800
+++ mmult_float.cpp.L1       2022-03-04 16:58:38.673526756 +0800
@@ -72,6 +72,7 @@
L1:
    for (int i = 0; i < BATCH; i++) {
        // Iterate over output classes
    +#pragma HLS PIPELINE II = 1
    L2:
        for (int j = 0; j < CLASSES; j++) {
            // Perform the dot product
```

Figure 3: Inserting HLS directive for L1.

Table 6: Loop details for L1 pipelining

| Loop Name    | Latency (cycles) |      | Iteration<br>Latency | Initiation Interval |        | Trip<br>Count | Pipelined |
|--------------|------------------|------|----------------------|---------------------|--------|---------------|-----------|
|              | min              | max  |                      | achieved            | target |               |           |
| - LOAD_OFF_1 | 5                | 5    | 1                    | -                   | -      | 5             | no        |
| - LOAD_W_1   | 1300             | 1300 | 130                  | -                   | -      | 10            | no        |
| + LOAD_W_2   | 128              | 128  | 1                    | -                   | -      | 128           | no        |
| - LOAD_I_1   | 2064             | 2064 | 258                  | -                   | -      | 8             | no        |
| + LOAD_I_2   | 256              | 256  | 2                    | -                   | -      | 128           | no        |
| - L1         | 1402             | 1402 | 1291                 | 16                  | 1      | 8             | yes       |
| - STORE_O_1  | 136              | 136  | 17                   | -                   | -      | 8             | no        |
| + STORE_O_2  | 15               | 15   | 3                    | -                   | -      | 5             | no        |

## 1.3 C. Increasing Pipeline Parallelism by Repartitioning Memories (8 marks)

Report

1. the design latency in cycles,
2. the overall device utilization (as Total per Resource),
3. the number of floating point adders and multipliers (you can find this information under the Instance section of the synthesis report) and
4. the Initiation Interval of the loops you pipelined.

<sup>1</sup>DSP (363%), FF (390%) , LUT (457%)

#### **1.4 D. Amortizing Iteration Latency with Batching (8 marks)**

Report

1. the design latency in cycles, and
2. the overall device utilization (as Total per Resource).

#### **1.5 E. Extending Batch Size with Tiling (8 marks)**

Report

1. the design latency in cycles, and
2. the overall device utilization (as Total per Resource).

#### **1.6 F. Hardware compilation and FPGA testing on the PYNQ (8 marks)**

Report

1. the measured speedup and
2. measured classification accuracy.

### **2 Part 2: Fixed-Point Optimizations (30 marks)**

1. the fixed-point validation accuracy reported by mnist.py after you've tweaked the SCALE factor.
2. the design latency in cycles
3. the overall device utilization (as Total per Resource).
4. your measured system speedup over the fixed-point CPU implementation
5. your measured classification accuracy on the 8k MNIST test sample
6. how many multipliers are instantiated in your desing?
7. report the initiation interval of the matrix multiplication loop that you pipelined
8. given the number of multipliers in your design and input throughput via the AXI port, is the design bandwidth- or compute-limited?

### **3 Part 3: Open-ended design optimization (30 marks)**

Vitis High-Level Synthesis User Guide